
THE SUPERHUMAN SELLER

The Superhuman Seller **Playbook**

How B2B reps are using AI to build pipeline, score deals honestly, and forecast a number they can actually defend. Written by the operator who built the engine behind it.

~12,700 lines of production Python

25+ automation scripts

15+ live cron jobs

10 systems wired together

Read this part first

I'm not going to tell you AI is the future of sales. You've heard that enough times that it stopped meaning anything. I'd rather show you what it actually does when you put it to work, and let you decide.

Here's where I'm coming from. Over the past few months I've built and shipped a working sales-automation stack: about 12,700 lines of production Python across more than 25 scripts, with 15-plus cron jobs running around the clock. It pulls call transcripts, scores deals against MEDDPICC, drafts outbound, watches reply inboxes every fifteen minutes, and writes research briefs into a CRM before a rep ever opens it. None of that is a demo. It runs on real pipeline, every day.

I'm staying anonymous on purpose. This isn't about me or my logo. It's about the method, and the method works regardless of whose name is on the door. What I'll give you in these pages is the same thinking that went into the engine, translated so a seller can run a lighter version of it with nothing more than a chat window and a bit of discipline.

One promise: no fluff about "synergy" or "transformation." Just the moves that move a number.

"The future is not fewer people. It's different work, new roles, and better rewards for those who embrace it."

ZEB JONES, CEO OF CLICKUP · MAY 2026

That line is the whole reason I wrote this. The work is changing. The reps who treat AI as a tool they direct, not a threat they ignore, are about to have a very good couple of years. Let's get you there.

CONTENTS

What's inside

- 01 The sorting: why builders and sellers are safe →
- 02 What "AI for sellers" actually means →
- 03 The operating loop behind a superhuman rep →
- 04 Pipeline generation: find, qualify, reach out fast →
- 05 Deal scoring with MEDDPICC (and how to automate it) →
- 06 Forecasting a number you can defend →
- 07 The 15 prompts →
- 08 Career insurance for the next three years →
- 09 The 90-day mastery plan →

The sorting

There's a quiet reshuffling happening inside companies right now, and most layoff memos are too polite to name it plainly. One CEO wasn't.

Cloudflare's Matthew Prince wrote that the "builders and sellers" were safe at his company, while AI was coming for the "measurers." He described cutting a marketing team that, "like in most companies, was teeming with measurers."

MATTHEW PRINCE, CEO OF CLOUDFLARE, IN THE WALL STREET JOURNAL · MAY 2026

Strip away how blunt that is and look at the logic. A measurer is someone whose job is to count, report, and route information that already exists. That's exactly the work software is best at absorbing. A builder makes something that didn't exist before. A seller creates revenue that wouldn't have happened otherwise. Those two are hard to automate because they require judgment, relationships, and the willingness to be wrong in public.

Sales sits on the safe side of that line. But there's a catch, and it's the whole point of this playbook: you only stay on the safe side if you do the part of the job that can't be automated, and you let the machine take the part that can. The rep who spends their morning copying notes between tabs, building lists by hand, and writing the same follow-up for the hundredth time is doing measurer work inside a seller's title. That's the person at risk. Not because they sell, but because they've quietly turned selling into clerical work.

The split inside your own week

Pull up last week and sort it into two buckets. In one bucket, the conversations: discovery calls, demos, negotiation, the email where you actually moved a deal. In the other, the prep and admin: research, list building, note-taking, CRM updates, status reports, drafting the seventh version of an outreach email. Most reps I've watched are spending sixty percent or more of their time in the second bucket.

The promise of doing this well isn't "work less." It's flipping that ratio. Give the prep and the admin to a system you trust, and you walk into every conversation more prepared than the rep across the table who did it by hand, while having more conversations than they could fit in a day. That's what "superhuman" means here. Not faster typing. More at-bats, each one sharper.

FROM THE ENGINE

The reason I trust this split is that I automated the second bucket and measured what came back. A **post-call brief pipeline** runs every four hours, pulls the transcript from whichever source caught the call (Zoom, Granola, or the CRM body), and writes a structured summary plus a full MEDDPICC read into the CRM before the rep has logged anything. That's prep and admin a human used to do after every single call. The rep now reviews instead of writes.

What "AI for sellers" actually means

Forget the marketing. A language model is good at exactly three things that matter to a seller, and useless at a fourth that everyone keeps asking it to do.

The three things it's genuinely good at

Reading fast and summarizing without getting bored. Hand it a 40-minute transcript, a 10-K, three news articles, and a LinkedIn profile, and it will pull out what matters in seconds without skimming the way a tired human does at 4pm. This is the single highest-leverage use for a rep, and almost nobody uses it enough.

Drafting in a voice and structure you specify. It will write the cold email, the follow-up, the recap, the mutual action plan. Not perfectly. But a good first draft you edit in ninety seconds beats a blank page you stare at for ten minutes.

Pattern-matching against a framework. Give it MEDDPICC, or your qualification criteria, or your ICP definition, and it will hold a messy conversation up against that structure and tell you what's missing. Humans do this too, but they forget a dimension, or they grade their own deals generously. The model doesn't have happy ears.

The thing it's bad at, that you must own

Judgment about people. Whether a champion is actually a champion or just polite. Whether silence means "no" or "busy." When to push and when to wait. Whether the deal is real. The model can surface signals and lay out options, but the call is yours, and it always will be. That's the part that keeps you employed, so guard it.

THE MENTAL MODEL

Treat the AI like a brilliant, tireless analyst who started yesterday. It reads everything, drafts anything, and never complains. But it has zero context you don't give it, and it will state a wrong answer with total confidence. Your job shifts from doing the work to **directing and checking** the work. That shift is the entire skill.

Why "garbage in" hits sellers harder

A model is only as good as the context you feed it. Ask "write me a cold email" with nothing else and you'll get something generic that sounds like every other AI email flooding inboxes right now. Feed it the

prospect's recent funding round, the role they just hired for, the specific pain that role implies, and your actual product's angle on it, and you get something a human would reply to. The skill isn't prompting. It's knowing what context matters, which happens to be the same instinct that made you good at sales in the first place.

FROM THE ENGINE

Every brief the system writes runs on assembled context, not vibes. The pre-call generator pulls CRM history, checks a prospect knowledge base, and runs four live web searches (news, funding, hiring, exec changes) before it writes a word. The research layer uses **18 validated research processes running at 90%+ accuracy**. The output is only good because the input is built, deliberately, every time.

SECTION 03

The operating loop

A superhuman rep isn't running fifty disconnected tricks. They're running one loop, over and over, with AI doing the heavy lifting at each turn. Learn the loop and every tactic in this book has a home.

STAGE	WHAT YOU DO	WHAT THE AI DOES
Find	Define who's worth your time. Set the ICP.	Scans signals (funding, hires, tech changes), scores accounts, builds the list.
Research	Pick the angle. Decide the play.	Reads everything, writes a brief, suggests the opener and likely objections.
Reach out	Approve, personalize the top line, hit send.	Drafts the sequence in your voice, tracks replies, flags the live ones.
Qualify	Run the conversation. Listen.	Transcribes, scores against MEDDPICC, names the gaps.
Advance	Decide the next step. Make the human calls.	Drafts the recap, the action plan, the multithread email.
Forecast	Commit a number you'll stand behind.	Weights the pipeline, writes the deal narratives, flags what's gone quiet.

Notice the pattern in the left column. Every single thing you do is a decision or a relationship. Everything in the right column is reading, writing, and remembering at scale. That division is not an accident. It's the design principle. When you find yourself doing something from the right column by hand, that's the signal to build or delegate.

25+	~12.7k	15+	10
SCRIPTS IN PRODUCTION	LINES OF PYTHON	CRON JOBS LIVE	TOOLS WIRED TOGETHER

The engine behind this brand connects HubSpot, Zoom, Granola, Smartlead, LeadIQ, Google Workspace, Linear, Slack, Clay, and a MEDDPICC scoring layer. You don't need all of that to start. You need the loop and one good chat window. The rest is just turning the loop up.

Pipeline generation

Pipeline is the part of the job most reps do worst, because it's tedious and there's always a reason to do it tomorrow. AI removes the tedium, which removes the excuse.

Stop building lists. Start scoring signals.

The old way to build a list: filter by industry and headcount, export, start dialing the top. The problem is that a company matching your ICP on paper might have zero reason to buy this quarter. A worse-fit company that just raised a round, hired a VP in your buyer's seat, or got named in a breach has a reason to buy right now. Timing beats fit more often than reps admit.

So flip it. Instead of asking "who matches my profile," ask "who matches my profile AND just did something that creates pain my product solves." That second clause is a signal, and signals are exactly what AI is good at finding and ranking across hundreds of accounts while you sleep.

FROM THE ENGINE

The signal monitor I built scores each account on a simple, tunable model: funding and exec changes weigh heaviest, recency matters (something from the last 30 days outranks something from last quarter), and ICP keywords like security, governance, and compliance add points. Anything that crosses the threshold gets a drafted outreach blurb and a Slack ping. A daily call-down list of **50 ICP-fit contacts** gets built and brief-stamped before the workday starts. The rep wakes up to a ranked queue, not a blank spreadsheet.

The three-pass outreach you can run by hand

You don't need a custom engine to apply the idea. Here's the lightweight version, runnable in any chat window:

1. **Signal pass.** Paste in a prospect's recent news, funding, hiring, or LinkedIn activity. Ask the model what pain each item likely creates for your specific buyer persona.
2. **Angle pass.** Pick the strongest pain. Ask the model to connect it to one concrete thing your product does — not a feature dump, one angle.
3. **Draft pass.** Have it write a short email in your voice built entirely around that angle, with one clear ask.

The whole sequence takes about four minutes per prospect once you're warmed up. The emails land differently because they're built on a reason, not a template. A generic AI email is worse than no email. A

specific one outperforms what most reps write by hand because it never skips the research step out of laziness.

A WARNING WORTH THE PAGE

Inboxes are filling with obvious AI outreach right now, and buyers have learned the smell of it. Volume without specificity is how you burn a domain and a reputation. The edge isn't sending more. It's sending fewer, sharper messages that prove you actually did the reading. Use the time AI gives you to go deeper, not just wider.

The prompts in Section 07 cover the research brief, the signal scan, the cold email, and the follow-up sequence. Those four alone will rebuild how your top of funnel feels.

Deal scoring with MEDDPICC

Most reps don't have a forecasting problem. They have a qualification problem that only shows up as a forecasting problem at the end of the quarter. MEDDPICC fixes it upstream, and AI makes it something you'll actually keep up.

The eight dimensions, fast

DIMENSION	THE QUESTION IT FORCES YOU TO ANSWER
Metrics	What measurable outcome does the buyer get, in their numbers, not yours?
Economic Buyer	Who can actually say yes and release the budget? Have you met them?
Decision Criteria	What boxes must be checked for them to choose anyone?
Decision Process	What are the literal steps and dates from here to signature?
Paper Process	Procurement, legal, security review — what slows the signature down?
Identify Pain	What breaks for them if they do nothing? How much does that cost?
Champion	Who sells for you when you're not in the room? Have you tested them?
Competition	Who else is in the deal, including "do nothing" and an internal build?

The framework isn't the insight. Every rep has seen it. The insight is that nobody fills it out honestly or consistently, because doing it well means reviewing your own call notes against eight dimensions after every conversation, and that's exactly the kind of work reps skip when they're busy. So the scorecard ends up half-empty, optimistic, and stale, which is worse than not having one at all.

Let the machine grade the call

This is where AI changes the framework from a chore into a habit. After a call, feed the transcript to the model with the eight dimensions and ask it to score each one, cite the exact quote that supports the score, and flag what's missing. It doesn't have happy ears. It won't mark "Economic Buyer: confirmed" because the champion said "I'm pretty sure my VP will love this." It'll tell you that you've never actually spoken to the person who signs.

FROM THE ENGINE

In the live system, every call gets a **full 8-dimension MEDDPICC assessment** generated automatically and written back to the deal record. Seven of the eight dimensions map to structured CRM fields with a completion score, so a manager can see exactly how qualified the pipeline is without nagging anyone for updates. Latest call wins — the score reflects reality now, not what was true three weeks ago. That's the difference between a scorecard and a museum piece.

You can run a manual version with Prompt 08 (MEDDPICC gap analysis) in the next section. Do it after every meaningful call for one month and you'll never look at your pipeline the same way. The deals you were sure about will reveal their holes, and the quiet deals will reveal their strength.

THE HONESTY DIVIDEND

A scorecard the AI fills in from the actual transcript is harder to lie to yourself with. That discomfort is the value. A deal that looks weak on paper in week two is a deal you can still save. A deal that looked great until it slipped on the last day of the quarter is just a lesson you paid full price for.

Forecasting a number you can defend

A forecast isn't a wish. It's a claim you're willing to repeat to your boss's boss. The reps who get trusted with bigger territories are the ones whose number comes true, and that's a discipline you can systematize.

Probability-weighted, not gut-weighted

The simplest honest forecast assigns each deal a real probability of closing in the period, multiplies it by the deal size, and sums the result. The trap is where the probability comes from. Most reps pull it from how they feel about the deal, which is how a \$200k deal they like beats a \$200k deal they're avoiding, regardless of which one is actually further along.

Tie the probability to qualification instead. A deal with a confirmed economic buyer, a written decision process, a tested champion, and known competition is genuinely more likely to close than one missing four of those, and it should carry a higher weight. This is why Section 05 comes before this one. Your MEDDPICC completeness *is* your probability input. When the scoring is automated and honest, the forecast stops being a feeling and becomes math.

A WORKED EXAMPLE

DEAL	SIZE	MEDDPICC STRENGTH	WEIGHT	WEIGHTED
Acme renewal expansion	\$120k	Strong (7/8 confirmed)	80%	\$96k
Northwind new logo	\$200k	Medium (no paper process)	45%	\$90k
Beacon platform deal	\$300k	Weak (no economic buyer)	20%	\$60k
			Forecast	\$246k

The \$300k deal you're excited about contributes less than the \$120k deal you've actually de-risked. That feels wrong to a rep running on optimism. It's exactly right to a rep who wants to be believed in March and still employed in April.

Where AI earns its keep here

Two jobs. First, it keeps the inputs current — pulling the latest qualification state off every deal so your weights reflect this week, not last month. Second, it writes the narrative. For each deal in your commit, have the model produce three sentences: what's true, what's the risk, what's the next step and date. A

forecast review where every deal has a tight, honest narrative is a forecast people trust. One where you wing it deal by deal is one they discount.

FROM THE ENGINE

A deal-stage monitor watches the pipeline for movement and for silence. New deal, stage change, or a deal gone dark for 14+ days each triggers the right research: a pre-meeting brief, a competitive intel pull, or a re-engagement angle. The forecast stays honest because nothing rots quietly. The deal that went dark gets surfaced while there's still time to act, not discovered during the end-of-quarter scramble.

Prompt 14 turns a messy pipeline export into a weighted forecast with per-deal narratives in one pass. It's the prompt I'd hand a rep first if I could only give them one.

SECTION 07

The 15 prompts

These are starting points, not magic words. Swap in your product, your buyer, your voice. The variables in **green** are the parts you replace. Edit the output — always. The draft gets you to ninety percent; the last ten percent is the part that sounds like you.

PROMPT 01

Pre-call research brief

Use before any first meeting

You're my sales research analyst. Build me a pre-call brief on **[COMPANY]** for a meeting with **[NAME, TITLE]**. I sell **[ONE-LINE PRODUCT]** to **[BUYER TYPE]**.

Cover: company snapshot, anything from the last 90 days that signals timing (funding, hires, launches, incidents), the likely pain my product speaks to, two openers tied to a real signal, and the three objections this persona is most likely to raise. Keep it to one page. Flag anything you're unsure about instead of guessing.

PROMPT 02

Account signal scan

Use to prioritize a list

Here is recent activity for **[COMPANY]**: **[PASTE NEWS / FUNDING / HIRES / POSTS]**.

For each item, tell me: what pain or initiative it implies for **[BUYER PERSONA]**, how time-sensitive it is (now / this quarter / background), and whether it's a reason to reach out this week. Rank the items. Be skeptical — tell me which ones are noise.

PROMPT 03

ICP fit check

Use to qualify in or out fast

My ICP is: **[INDUSTRY, SIZE, TRIGGERS, DEAL-BREAKERS]**. Here's an account: **[PASTE DETAILS]**.

Score the fit 1–10 against my ICP, list what matches and what doesn't, and tell me the single biggest reason to pursue or pass. If it's a pass, say so plainly so I don't waste a week on it.

PROMPT 04

Cold email, built on a signal

Use for first touch

Write a cold email to **[NAME, TITLE]** at **[COMPANY]**. The reason I'm reaching out now: **[SPECIFIC SIGNAL]**. What I help with: **[ANGLE, NOT FEATURE LIST]**.

Under 90 words. No "I hope this finds you well." Open with the signal, not with me. One clear ask. Sound like a human who did their homework, not a template. Give me two subject lines.

PROMPT 05

Follow-up after no reply

Use 3–4 days after first touch

No reply to my first email (pasted below). Write a follow-up that adds a new angle or a piece of value – not "just bumping this." Shorter than the first. Make it easy to say yes to a 15-minute call, and easy to say "not now" without guilt.

First email: **[PASTE]**

PROMPT 06

The honest breakup

Use to close a dead thread

Write a short, low-pressure final email to **[NAME]** after several unanswered touches. Make it genuinely easy to walk away, name that I'll stop reaching out, and leave one door open. No guilt-tripping. These get more replies than people expect – keep it warm and brief.

PROMPT 07

Discovery question set

Use to prep a first call

I'm meeting **[NAME, TITLE]** at **[COMPANY]** about **[SUSPECTED PAIN]**. Give me 8 open discovery questions that uncover metrics, decision process, and pain – in plain language, not interrogation mode. Order them so the call feels like a conversation. Flag the two that matter most so I don't lose them if we run short.

PROMPT 08

MEDDPICC gap analysis

Use after every real call

Here's my call transcript/notes: **[PASTE]**.

Score this deal on all 8 MEDDPICC dimensions (Metrics, Economic Buyer, Decision Criteria, Decision Process, Paper Process, Identify Pain, Champion, Competition). For each: a 0–2 score, the exact quote that supports it, and what's missing. End with the single most important thing I don't yet know. Don't be generous.

PROMPT 09

Objection handling, live-ready

Use to prep or recover

The objection: **[PASTE WHAT THEY SAID]**. Context: **[DEAL STAGE, WHAT THEY SELL, WHO SAID IT]**.

Give me three ways to respond: one that acknowledges and reframes, one that asks a question to get underneath it, and one that uses a relevant proof point. Keep each to a couple of sentences I could actually say out loud.

PROMPT 10

Competitive positioning

Use when a rival enters the deal

I'm up against **[COMPETITOR]** at **[COMPANY]**. They care most about **[DECISION CRITERIA]**. Without trashing the competitor, help me position around where I'm genuinely stronger, and prep me for how they'll position against me. Tell me the one trap to avoid in this matchup.

PROMPT 11

Multithread / champion enablement

Use to widen a single-threaded deal

My only contact at **[COMPANY]** is **[NAME, TITLE]**. I need to reach **[OTHER ROLE]** too. Write a short message my champion can forward internally that makes them look good for sharing it, plus a separate note I'd send the new stakeholder directly. Frame both around outcomes, not my product.

PROMPT 12

Post-call recap + next steps

Use within an hour of a call

Turn these notes into a recap email to [NAME]: [PASTE NOTES].

Confirm what we agreed, restate their goal in their words, and lay out clear next steps with owners and dates. Professional, warm, short. End with one specific scheduling ask, not "let me know."

PROMPT 13

Mutual action plan

Use to formalize a late-stage deal

Build a mutual action plan to get [COMPANY] from where we are ([CURRENT STEP]) to signature by [TARGET DATE]. List each step, who owns it (them or me), and the date. Include the easy-to-forget ones: security review, legal, procurement. Format it so I can paste it into a shared doc.

PROMPT 14

Weighted forecast + narratives

Use before every forecast call

Here's my pipeline for the period: [PASTE DEALS WITH SIZE + STAGE + MEDDPIC NOTES].

Assign each deal a close probability based on qualification strength (not optimism), show the weighted value, and total the forecast. For each deal write three sentences: what's true, the biggest risk, the next step and date. Flag any deal I'm probably overcalling.

PROMPT 15

LinkedIn touch that isn't cringe

Use to warm a cold account

Write a short, no-pitch LinkedIn note to [NAME, TITLE] tied to [SOMETHING REAL THEY POSTED OR DID]. Sound like a peer, not a vendor. No "I'd love to pick your brain." The goal is one genuine reply, not a meeting or message one.

Career insurance

Let's go back to where we started. AI is sorting roles, and sales is on the protected side — if you do the protected work. Here's how to make sure you're on the right side of that line for the next three years.

"We're already seeing entirely new roles emerge, like Agent Managers, that didn't exist a year ago."

ZEB JONES, CEO OF CLICKUP · MAY 2026

Read that as a job posting, not a headline. The reps who win the next cycle are the ones who can direct a set of AI tools the way a manager directs a team: assigning the right work, checking the output, and owning the result. You don't need to write 12,700 lines of Python. You need to think like the person who would. That mindset — "what here is a decision, and what here is just reading and writing I can hand off" — is the durable skill.

Three habits that compound

Become the person who closes, not the person who reports. Wix's CEO described flattening his company so there's "less distance between the people setting direction and the people building the product." In a flatter org, the people closest to revenue get more trust and less oversight. Be unmistakably close to revenue.

Build a personal stack. You don't need engineering. You need three or four prompts you run every day until they're muscle memory, plus whatever tools your company gives you, wired into a routine. The rep with a system beats the rep with talent and no system, every quarter, predictably.

Get visibly good at the human part. When the research, drafting, and admin are handled, what's left is judgment, relationships, and nerve. Those have always been what separated great reps from average ones. AI doesn't change that. It just strips away the busywork that used to hide it.

THE UNCOMFORTABLE VERSION

If your honest answer to "what do I do that AI can't" is "I update the CRM and send follow-ups," that's not a moat. It's a notice. The good news is that the fix is the same thing that makes you better at the job today: hand off the measurer work, and spend the recovered hours being a better seller.

The 90-day mastery plan

You can't adopt all of this at once, and trying is how people quit in week two. Here's the order I'd run it in. One change at a time, until each one is a habit you don't think about.

MONTH 01 · RESEARCH & REPS

Make every conversation prepared

Pick two prompts and run them daily until they're automatic: the pre-call research brief (01) and the MEDDPICC gap analysis (08). Before every first meeting, generate the brief. After every real call, score the deal. Don't add anything else this month. The goal is simple: you walk into rooms more prepared than anyone across the table, and you stop lying to yourself about deal health. By week four this should feel like brushing your teeth.

MONTH 02 · TOP OF FUNNEL

Rebuild how you prospect

Now add the front end. Run the signal scan (02) and ICP check (03) on your target accounts every Monday, and let them rank your week. Replace your templated outreach with the signal-built cold email (04) and follow-up (05). Track one number: reply rate. It should climb, because every message now has a reason behind it. Fewer, sharper touches. This is the month your pipeline starts to feel different.

MONTH 03 · FORECAST & SYSTEM

Become the rep whose number is true

Bring it together. Run the weighted forecast (14) before every forecast call, with a tight narrative on each committed deal. Use the mutual action plan (13) on every late-stage deal. By now your MEDDPICC scoring from Month 1 feeds your forecast honestly, so your number is math, not mood. The outcome you're after: your manager stops double-checking your commit, because it keeps coming true. That reputation is worth more than any single deal.

After 90 days

You'll have a loop that runs whether you feel like it or not, and a stack of prompts that have become instinct. That's the foundation. From there it's just turning the dial up — more accounts, deeper research, tighter forecasting — and eventually, if you want, wiring pieces of it together so they run without you, the

way the engine behind this brand does. But that's a later chapter. Ninety days of the basics will already put you ahead of most of the floor.

This is the jumping-off point, not the whole journey.

I'll keep publishing the deeper builds — the actual workflows, the scoring models, the outbound architecture — in the newsletter, broken down so a seller can use them without an engineering team. If this playbook earned a few minutes of your attention, the next pieces are where it gets specific.

Subscribe at thesuperhumanseller.com

Written by the operator behind the engine. No name on the door — just the method, and the receipts.